

One prompt. Two measurement channels.

Independent journal-conditioning calibration · fixed model weights · no agent interaction

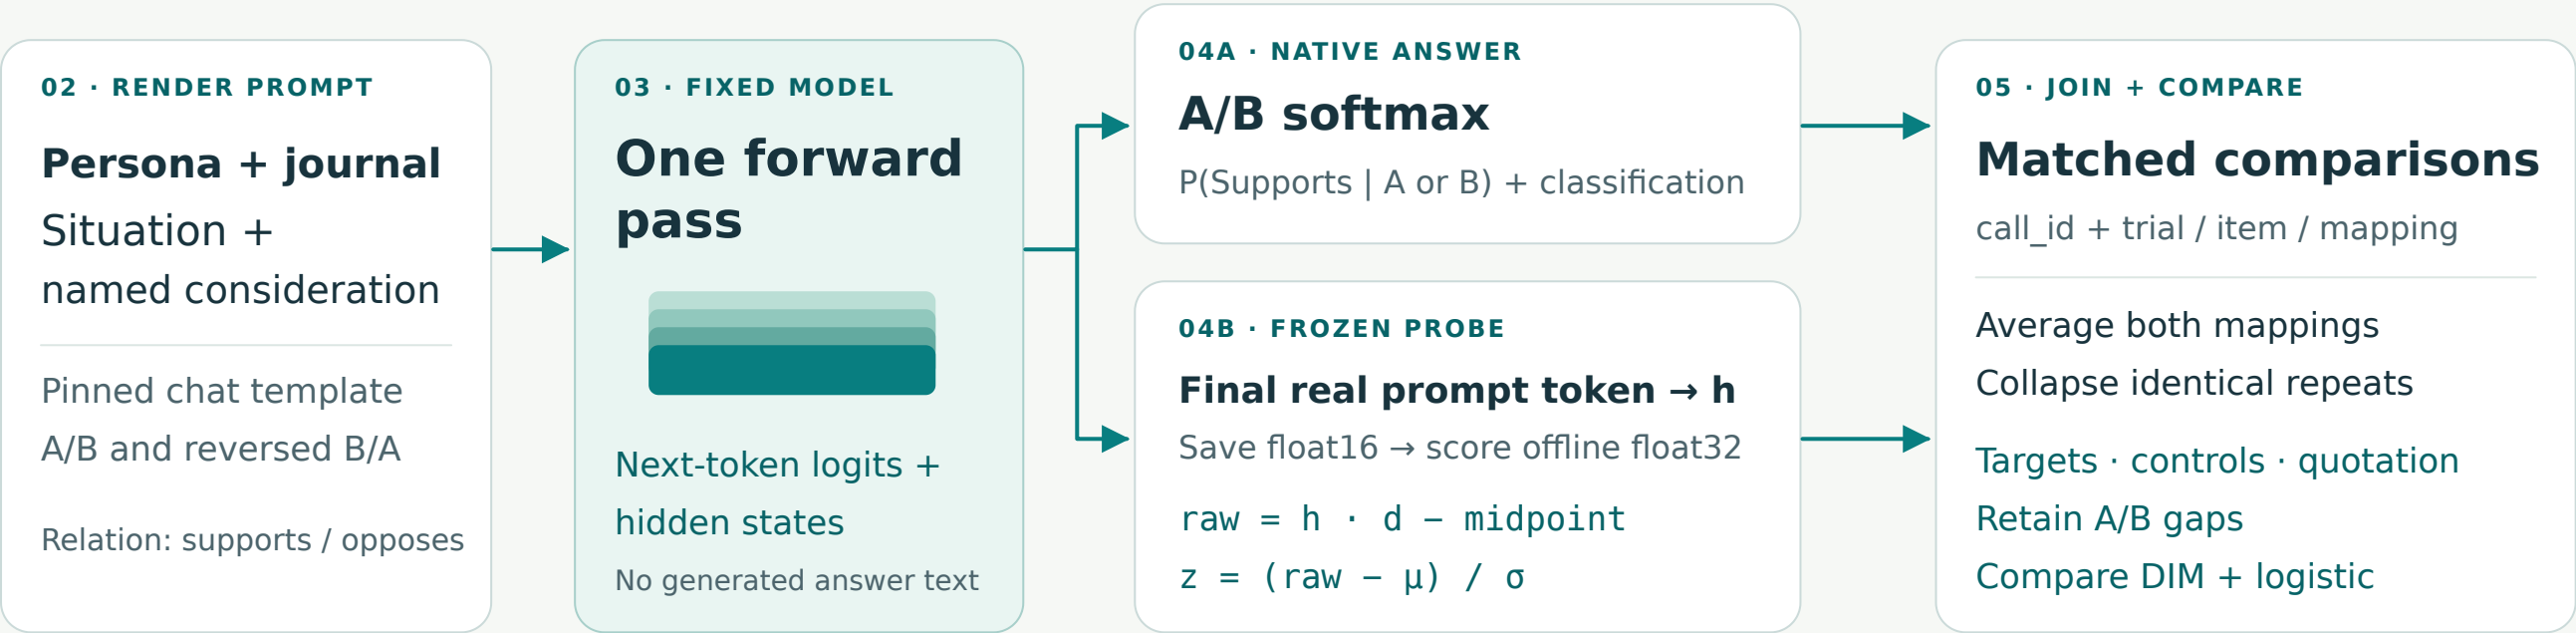
01 · MATCHED DESIGN · PER MODEL · 8 AUTHORED ANCHORS (4 HONESTY, 2 FAIRNESS, 2 SAFETY)

3 personas × 3 repeats × 8 anchors × 4 contexts = 288 observations

Neutral · Identical repeat · Seeded conviction · Quoted without adoption

Both A/B mappings → 576 calls/model

EACH MAPPING IS ONE FORWARD EVALUATION



Frozen weights + scalers; semantic Supports direction is unchanged by an A/B swap.

EXTRACTION + NUMERICAL GATE

Pinned revisions · BF16 · sequential A6000 (48 GB) runs

Gemma 2 9B IT

HF 28 = block 27

3,584 dimensions

Batch-two gate: pass → batch 2

Llama 3.1 8B Instruct

HF 20 = block 19

4,096 dimensions

Batch-two gate: fail → batch 1

48 numerical-check requests/model. The parallel two-GPU launcher was not used for the completed calibrations.

PROPOSED FUTURE STUDY · NOT RUN IN THIS PILOT

Seeded agent → Shared board → Recipient journal → Same battery

Validate measurement first; then matched seeded/unseeded recipients, pre-seeding baseline and persistence checks.

Supported: context-sensitive readouts. Moral adoption, persistence and contagion have not been established.