

When context moves a probe

Completed calibration · 6 September 2026 · teammate brief

Gemma’s categorical answers stayed stable while its continuous probe scores moved. This demonstrates context sensitivity. Moral adoption, persistence, and agent-to-agent contagion have not been established.

Before testing social influence, distinguish relation recognition from personal commitment. [Value Kaleidoscope](#) (Sorensen et al., 2024) motivates Supports/Opposes relations; [Geometry of Truth](#) (Marks & Tegmark, 2024) motivates linear readouts. Neither establishes an alignment measure.

EXACT METHODS

3 personas × 3 identical repetitions × 8 items × 4 journals = 288 observations/model; 2 A/B mappings = 576 captured vectors/model. Journals were neutral, identical neutral repeat, asserted honesty-absolutist conviction, and that conviction quoted without adoption. Four Honesty targets and four Fairness/Safety controls were authored examples. These were independent requests: no agents interacted.

Fixed checkpoints used unquantized BF16 inference on an A6000, without training or refitting. Capture: final non-padding prompt token, before the answer (HF tuple includes embeddings).

Checkpoint	HF index / width	Model and tokenizer revision
Gemma 2 9B IT	28 / 3,584	11c9b309abf73637e4b6f9a3fa1e92e615547819
Llama 3.1 8B Instruct	20 / 4,096	0e9e39f249a16976918f6564b8830bc894c89659

Native scores averaged next-token conditional A/B probabilities mapped to Supports. Float16 captures were scored in float32 with [Jeff’s frozen difference-in-means \(DIM\) probe](#): **s = h·direction – midpoint; z = (s – frozen mean)/frozen SD**. Swapped symbols never negate the probe. Logistic was secondary. Gemma used batch 2; Llama used batch 1 after failing the batch-2 numerical gate.

COMPLETED RESULTS

Gemma matched every reference, with mapping agreement **288/288** and no classification flips. Llama mapping agreement was **144/288**. Both yielded **576/576 finite vectors**. Frozen [evidence](#) gives:

Mean absolute DIM change	Gemma: Honesty	Gemma: controls	Llama: Honesty	Llama: controls
Identical repeat – neutral	0.000	0.000	0.000	0.000
Seeded – neutral	0.158	0.099	0.119	0.093
Quoted – neutral	0.119	0.122	0.195	0.090

Units: each model’s frozen selection-score SD, not a new-battery effect size. Each cell describes **12 persona-item pairs**, after collapsing repeats and averaging mappings; these are not independent simulation replications. Mean absolute A/B probe gaps over all items/contexts were **0.124 (Gemma)** and **1.304 (Llama)**.

INTERPRETATION AND NEXT STEP

Quotation and controls also moved scores. DIM and logistic disagreed on the sign of reference-oriented seeded target changes. Absolute movement therefore cannot establish stronger endorsement. Eight easy items, three personas, and one payload limit transfer claims. [Probe control tasks](#) (Hewitt & Liang, 2019) motivate checking alternative explanations.

Next: validate new situation families, paraphrases, and mappings; measure value trade-offs separately; distinguish exposure from stance and test persistence. Matched social simulations come after measurement validation.